




# GLMM - Bayesian Inference with Stan

## Introduction to Stan

L. Egidi - DEAMS, Units ([legidi@units.it](mailto:legidi@units.it))

PhD Course in Statistics - XXXV cycle

# Origins



**Stan**islaw Ulam (1909-1984): Manhattan project, H-Bomb experiments in Los Alamos, MCMC father jointly with John von Neumann.

# Indice

- 1 What is Stan?
- 2 Why Stan?
- 3 Writing a Stan program
- 4 Linked package: bayesplot

# What is Stan?

- **Probabilistic programming** language and inference algorithms.
- Stan **program**
  - declares data and (constrained) parameter variables
  - defines log posterior (or penalized likelihood)
- Stan **inference**
  - MCMC for full Bayes
  - Variational Bayes for approximate Bayes
  - Optimization for (penalized) MLE
- Stan **ecosystem**
  - lang, math library (C++)
  - interfaces and tools (R, Python, Julia, many more)
  - documentation (example model repo, [user guide & reference manual](#),  
[case studies](#), R package vignettes)
  - online community ([Stan Forums](#) on Discourse)

# Why Stan?


- Fit rich Bayesian statistical models. Close to the *big data* philosophy.
- Efficiency
  - Hamiltonian Monte Carlo + NUTS
  - Compiled to C++
- Flexible domain specific language
- “Freedom-respecting, open-source”
  - doc & written materials
  - interacting community
  - continuous development
- Interaction with some other R packages designed to explore the Stan output.

# Who is using Stan?

- **Biological & physical sciences:** clinical trials, epidemiology, genomics, population ecology, entomology, ophthalmology, neurology, agriculture, fisheries, cancer biology, astrophysics & cosmology, molecular biology, oceanography, climatology.
- **Social sciences:** population dynamics, psycholinguistics, social networks, political science, human development, economics.
- **Many more:** sports analytics, public health, publishing, finance, pharma, actuarial, recommender systems, educational testing, materials engineering.

# Interfaces

## Interfaces + Tools



# Indice

- 1 What is Stan?
- 2 Why Stan?**
- 3 Writing a Stan program
- 4 Linked package: bayesplot



# Improving MCMC performance

With Stan, we aim to provide an MCMC implementation that works robustly for as many target distributions as possible

- Gibbs, RW Metropolis can be very inefficient, hard to diagnose.
- To explore complicated high-dimensional spaces we need to leverage what we know about the geometry of the **typical set**.
- For such a reason, Stan enjoys **Hamiltonian Monte Carlo**.

The Stan users may use, analyze and interpret HMC outputs as they were *standard* MCMC outputs.

# Index

- 1 What is Stan?
- 2 Why Stan?
- 3 Writing a Stan program
- 4 Linked package: bayesplot

# Before starting

What is a Bayesian model?

- Building a Bayesian model forces us to build a model for how the data is generated
- We often think of this as specifying a prior and a likelihood, as if these are two separate things
- They are not!

# Generative models

The philosophy behind Stan is to think **generatively**.



The model is expressed as a joint probability distribution of observed and unobserved variables, which may be decomposed as follows:

$$p(y, \theta) = p(y|\theta)\pi(\theta) \quad (1)$$

The posterior of interest is then proportional to the joint distribution (1):

$$p(\theta|y) \propto p(y|\theta)\pi(\theta) \quad (2)$$

# Generative models

A Bayesian modeller commits to an a priori joint distribution:

$$p(y, \theta) = \underbrace{p(y|\theta)\pi(\theta)}_{\text{Likelihood} \times \text{Prior}} = \underbrace{\pi(\theta|y)p(y)}_{\text{Posterior} \times \text{Marginal Likelihood}} \quad (3)$$

# Generative models and vague priors

What is the problem with *vague/diffuse* priors?

- If we use an improper prior, then we do not specify a joint model for our data and parameters.
- More importantly, we do not specify a data generating mechanism  $p(y)$ .
- By construction, these priors **do not regularize inferences**, which is quite often a bad idea
- Proper but diffuse is better than .improper but is still often problematic.

# Generative models

- If we disallow improper priors, then Bayesian modeling is generative.
- In particular, we have a simple way to simulate from  $p(y)$ :

$$\begin{array}{ccc}
 \theta^* \sim \pi(\theta) & & \\
 \downarrow & \longleftrightarrow & y^* \sim p(y) \\
 y^* \sim p(y|\theta^*) & & 
 \end{array}$$

# Stan computations

Stan works in logarithmic terms: all the computations are actually done on log-scale. So, for the posterior we have.

$$\log(\pi(\theta|y)) = \log(\pi(\theta)) + \log(p(y|\theta)) + \text{constant} \quad (4)$$

Products become sums of logs:

$$p(y|\theta) = \prod_{i=1}^n p(y_i|\theta) \rightarrow \log(p(y|\theta)) = \sum_{i=1}^n \log(p(y_i|\theta)).$$



# Starting point

We are now going to write a Stan program together:

- Open a new empty file in RStudio
- Save it as `linear_regression.stan`

# Blocks strategy

Stan programs are organized into **blocks**:

- **data** block: declare data types, sizes, and constraints. Read from data source and constraints validated. Evaluated: once.
- **parameters** block: declare parameter types, sizes, and constraints. Evaluated: every log prob evaluation.
- **transformed parameters** block: declare those parameters transformed from the original ones declared in the parameters block. Evaluated: every log prob evaluation.
- **model** block: statements defining the posterior density in log scale. Evaluated: every log prob evaluation.
- **generated quantities**: declare and define derived variables. (P)RNGs, predictions, event probabilities, decision making. Constraints validated. Evaluated: once per draw.

# Data block

```
data {  
  // Dimensions  
  int<lower=1> N;  
  int<lower=1> K;  
  
  // Variables  
  matrix[N, K] X;  
  vector[N] y;  
  
}
```

```
// single line comment  
/* multiple lines of  
comments */
```

# Parameters' block

```
parameters {  
  real alpha;  
  vector[K] beta;  
  real<lower=0> sigma;  
}
```

constraints *required* in  
parameters block

# Model block

```
model {
  // priors (flat, uniform, if omitted)
  sigma ~ exponential(1);
  alpha ~ normal(0, 10);
  for (k in 1:K) beta[k] ~ normal(0, 5);

  for (n in 1:N) {
    y[n] ~ normal(X[n, ] * beta + alpha, sigma);
  }
}
```

Why is the default automatically uniform?

- $\pi(\theta) \propto 1$  (0 on log scale)
- Nothing added to log prob

## Generated quantities block

```
generated quantities {  
  vector[N] y_rep;  
  for (n in 1:N) {  
    real y_hat = X[n,] * beta + alpha; // local/temp  
    y_rep[n] = normal_rng(y_hat, sigma);  
  }  
}
```

# Complete Stan model

```

data {
  int<lower=1> N;
  int<lower=1> K;
  matrix[N, K] X;
  vector[N] y;
}

parameters {
  real alpha;
  vector[K] beta;
  real<lower=0> sigma;
}

model {
  sigma ~ exponential(1);
  alpha ~ normal(0, 10);
  for (k in 1:K) beta[k] ~ normal(0, 5);

  for (n in 1:N)
    y[n] ~ normal(alpha + X[n, ] * beta, sigma);
}

generated quantities {
  vector[N] y_rep;
  for (n in 1:N)
    y_rep[n] = normal_rng(alpha + X[n, ] * beta, sigma);
}

```

Observed  
variables

Unobserved  
variables

$\log \pi(\theta)$   
+  
 $\log p(y | \theta)$

Simulate from  
generative model

## Launching the Stan model from R

Now we may launch the Stan program directly in R:

```
library(rstan)

# passing the data (already stored)
data <- list(N=N, K=K, X=X, y=y)

# fitting the model
fit1 <- stan(
  file = 'linear_regression.stan',
  data = data,
  iter = 2000,
  chains = 4)

# extracting the estimates
sims <- extract(fit1)
```



## First example: 8 schools

This example studied coaching effects from eight schools.

We denote with  $y_{ij}$  the result of the  $i$ -th test in the  $j$ -th school. We assume the following model:

$$y_{ij} \sim \mathcal{N}(\theta_j, \sigma_y^2)$$

$$\theta_j \sim \mathcal{N}(\mu, \tau^2)$$

Do some schools perform better/worse according to these coaching effects?

Here is the data, already aggregated by schools:

```
schools_dat <- list(J = 8,
  y = c(28, 8, -3, 7, -1, 1, 18, 12),
  sigma = c(15, 10, 16, 11, 9, 11, 10, 18))
```

## First Stan model: 8 schools

```
// saved as 8schools.stan
data {
  int<lower=0> J;          // number of schools
  real y[J];              // estimated treatment effects
  real<lower=0> sigma[J]; // standard error of effect estimates
}
parameters {
  real mu;                // population treatment effect
  real<lower=0> tau;       // standard deviation in treatment effects
  vector[J] eta;          // unscaled deviation from mu by school
}
transformed parameters {
  vector[J] theta = mu + tau * eta; // school treatment effects
}
model {
  eta ~ normal(0,1);          // prior
  y ~ normal(theta, sigma);   //likelihood
}
```

## First example: 8 schools

To fit the model and visualize the estimates, it is sufficient to type in R the following commands (with 2000 iterations and 4 chains as a default):

```
fit_8schools <- stan(file = '8schools.stan', data = schools_dat)
print(fit_8schools, pars=c("mu", "tau", "theta"))
```

	mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
mu	7.89	5.04	-2.31	4.74	7.92	11.05	18.05	2352	1
tau	6.70	5.71	0.24	2.61	5.43	9.19	21.16	1480	1
theta[1]	11.36	8.23	-2.25	6.18	10.29	15.46	31.15	3161	1
theta[2]	7.89	6.21	-4.43	3.96	7.83	11.78	20.47	4923	1
theta[3]	6.05	7.59	-10.81	1.92	6.56	10.81	20.25	4057	1
theta[4]	7.60	6.44	-5.36	3.74	7.63	11.73	20.57	5055	1
theta[5]	5.13	6.23	-8.45	1.35	5.60	9.26	16.37	4346	1
theta[6]	5.95	6.68	-8.21	1.99	6.30	10.21	18.21	4313	1
theta[7]	10.62	6.93	-1.58	6.12	10.14	14.53	25.63	3381	1
theta[8]	8.40	7.77	-7.18	3.84	8.26	12.63	25.78	3854	1

# Indice

- 1 What is Stan?
- 2 Why Stan?
- 3 Writing a Stan program
- 4 Linked package: bayesplot**

# Posterior graphical analysis with bayesplot

Once we fit a model, it is to vital check it via graphical inspection. The bayesplot package (for any help, see the [vignette](#)) is designed to this task.



The package allows to display:

- Posterior uncertainty intervals
- Univariate marginal posterior distributions
- Bivariate plots
- Trace plots
- Posterior predictive plots

# Posterior graphical analysis with bayesplot


The first step is to save the posterior. Then you have many choices:

```
library(bayesplot)
posterior <- as.array(fit_8schools)


mcmc_intervals(posterior)      # posterior intervals
mcmc_areas(posterior)         # posterior areas
mcmc_dens(posterior)          # marginal posteriors
mcmc_pairs(posterior)         # bivariate plots
mcmc_trace(posterior)         # trace plots
```

With the arguments `pars` or `regex_pars` you may select the desired parameters.

# Posterior uncertainty intervals




# Posterior uncertainty areas







# Marginal posteriors



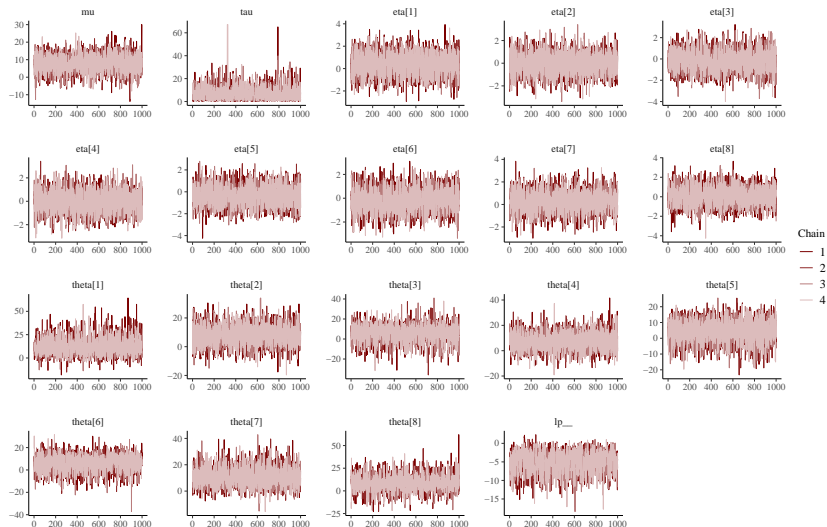
# Marginal posteriors separated for each chain



# Bivariate posterior plots



# Trace plots for the Markov chains



# Our challenge with Stan

The Stan shuttle is ready to start! We will learn to:

- **write** simple and more complex model in Stan: lm, glm, hierarchical models.
- **analyze** the posterior summaries.
- **criticize** the model and, eventually, change/reparametrize it.

## Further reading

### Further reading:

- Carpenter, B, and Gelman, A, Hoffman, M.D., Lee, D., Goodrich, B., Betancourt, M., Brubaker, M., Guo, J., Li, P., Riddell, A. (2017). Stan: A Probabilistic Programming Language, *Journal of statistical software* 76(1). Here the [pdf](#)

### Further optional reading about Hamiltonian Monte Carlo:

- Betancourt, M. (2017) A conceptual introduction to Hamiltonian Monte Carlo. Here the [pdf](#)